

data_exploration

October 15, 2023

1 Data Exploration

```
[243]: import numpy as np
import pandas as pd
import seaborn as sns
sns.set_theme()
```

1.1 Load data

```
[244]: df = pd.read_csv("data/deliverytime.csv")

print(df.columns)

df.describe()
```

```
Index(['ID', 'Delivery_person_ID', 'Delivery_person_Age',
      'Delivery_person_Ratings', 'Restaurant_latitude',
      'Restaurant_longitude', 'Delivery_location_latitude',
      'Delivery_location_longitude', 'Type_of_order', 'Type_of_vehicle',
      'Time_taken(min)'],
      dtype='object')
```

```
[244]:
```

	Delivery_person_Age	Delivery_person_Ratings	Restaurant_latitude \
count	45593.000000	45593.000000	45593.000000
mean	29.544075	4.632367	17.017729
std	5.696793	0.327708	8.185109
min	15.000000	1.000000	-30.905562
25%	25.000000	4.600000	12.933284
50%	29.000000	4.700000	18.546947
75%	34.000000	4.800000	22.728163
max	50.000000	6.000000	30.914057

	Restaurant_longitude	Delivery_location_latitude \
count	45593.000000	45593.000000
mean	70.231332	17.465186
std	22.883647	7.335122
min	-88.366217	0.010000

25%	73.170000	12.988453
50%	75.898497	18.633934
75%	78.044095	22.785049
max	88.433452	31.054057

	Delivery_location_longitude	Time_taken(min)
count	45593.000000	45593.000000
mean	70.845702	26.294607
std	21.118812	9.383806
min	0.010000	10.000000
25%	73.280000	19.000000
50%	76.002574	26.000000
75%	78.107044	32.000000
max	88.563452	54.000000

Some of the column names aren't great (e.g. "Time_taken(min)"). Let's rename them first.

```
[245]: cols = df.columns.tolist()
cols[-1] = "time_taken"
cols = [col.lower() for col in cols]
df.columns = cols
print(cols)
```

```
['id', 'delivery_person_id', 'delivery_person_age', 'delivery_person_ratings',
'restaurant_latitude', 'restaurant_longitude', 'delivery_location_latitude',
'delivery_location_longitude', 'type_of_order', 'type_of_vehicle', 'time_taken']
```

1.2 Missing values

```
[246]: df.isnull().sum()
```

```
[246]: id          0
delivery_person_id  0
delivery_person_age  0
delivery_person_ratings  0
restaurant_latitude  0
restaurant_longitude  0
delivery_location_latitude  0
delivery_location_longitude  0
type_of_order        0
type_of_vehicle       0
time_taken           0
dtype: int64
```

There are no missing values present in the dataset, so there is no need to impute or to drop any columns or rows.


1.3 Dependent variable

```
[247]: avg_time = df.time_taken.mean()
median_time = df.time_taken.median()
print(f"Average delivery time: {avg_time}")
print(f"Median delivery time: {median_time}")
hist = sns.histplot(df.time_taken)
hist.vlines(avg_time, 0, 2250, color = "black", label = "Avg. Time")
hist.vlines(median_time, 0, 2250, color = "green", label = "Median Time")
hist.set(xlabel = "Time taken")
hist.legend()
```

Average delivery time: 26.29460662821047

Median delivery time: 26.0

[247]: <matplotlib.legend.Legend at 0x7f9692638370>



The distribution has a positive skew, and the mean is slightly above the median. The fastest deliveries take only ten minutes, the slowest 54 minutes.

1.4 Categorical independent variables

```
[248]: cat_cols = df.select_dtypes(include='object')
cat_cols.head()
```

```
[248]:      id delivery_person_id type_of_order type_of_vehicle
0  4607      INDORES13DEL02      Snack      motorcycle
1  B379      BANGRES18DEL02      Snack      scooter
2  5D6D      BANGRES19DEL01      Drinks      motorcycle
3  7A6A      COIMBRES13DEL02      Buffet      motorcycle
4  70A2      CHENRES12DEL01      Snack      scooter
```

There are four categorical variables: delivery id, delivery person ID, the type of order and the type of vehicle. The id is largely meaningless, but it does seem to contain a restaurant ID.

```
[249]: import re

pattern = r"^.+RES\d{2}"

df["restaurant"] = [re.search(pattern, cell).group() for cell in df.
    ↪delivery_person_id]

restaurants = list(set(df.restaurant))
restaurants.sort()
print(len(restaurants))

df.sort_values(by=['restaurant']).loc[0:100, ["restaurant",
    ↪"restaurant_latitude", "restaurant_longitude"]]
```

418

```
[249]:      restaurant restaurant_latitude restaurant_longitude
0      INDORES13      22.745049      75.892471
11859  INDORES13      22.745049      75.892471
12357  INDORES13      22.745049      75.892471
14807  INDORES13      22.745049      75.892471
12346  INDORES13      22.745049      75.892471
...      ...      ...      ...
17549  PUNERES17      18.530963      73.828972
3759   PUNERES17      18.530963      73.828972
34206  PUNERES17      18.530963      73.828972
23936  PUNERES17      18.530963      73.828972
100    PUNERES17      18.530963      73.828972
```

[16433 rows x 3 columns]

As cases with the same restaurant ID share the same location, we can infer that there are a total of 418 restaurants.

```
[250]: cat_cols.drop("id", axis = 1, inplace = True)
```

1.4.1 Type of order

```
[251]: print(f"Number of unique order types: {cat_cols.type_of_order.nunique()}")
cat_cols.type_of_order.unique().tolist()
```

Number of unique order types: 4

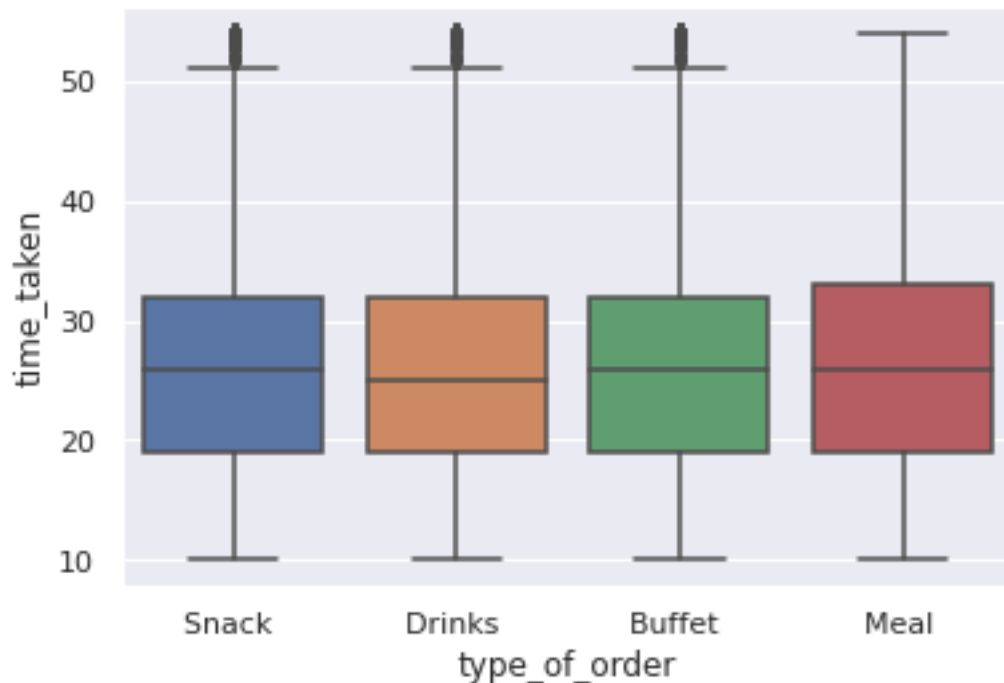
```
[251]: ['Snack ', 'Drinks ', 'Buffet ', 'Meal ']
```

The categories have superfluous whitespaces at the end of the string. These are removed first.

```
[252]: cat_cols["type_of_order"] = [cell.strip() for cell in cat_cols.type_of_order]
```

```
[253]: sns.boxplot(x = df.type_of_order, y = df.time_taken)
```

```
[253]: <AxesSubplot:xlabel='type_of_order', ylabel='time_taken'>
```



```
[254]: df.groupby("type_of_order")["time_taken"].describe()
```

```
[254]:
```

	count	mean	std	min	25%	50%	75%	max
type_of_order								
Buffet	11280.0	26.283511	9.411344	10.0	19.0	26.0	32.0	54.0

Drinks	11322.0	26.187953	9.298465	10.0	19.0	25.0	32.0	54.0
Meal	11458.0	26.419270	9.424849	10.0	19.0	26.0	33.0	54.0
Snack	11533.0	26.286309	9.399147	10.0	19.0	26.0	32.0	54.0

Meals tend to take slightly longer to deliver, but the difference isn't large.

1.4.2 Type of vehicle

```
[255]: print(f"Number of unique vehicle types: {cat_cols.type_of_vehicle.nunique()}")
cat_cols.type_of_vehicle.unique().tolist()
```

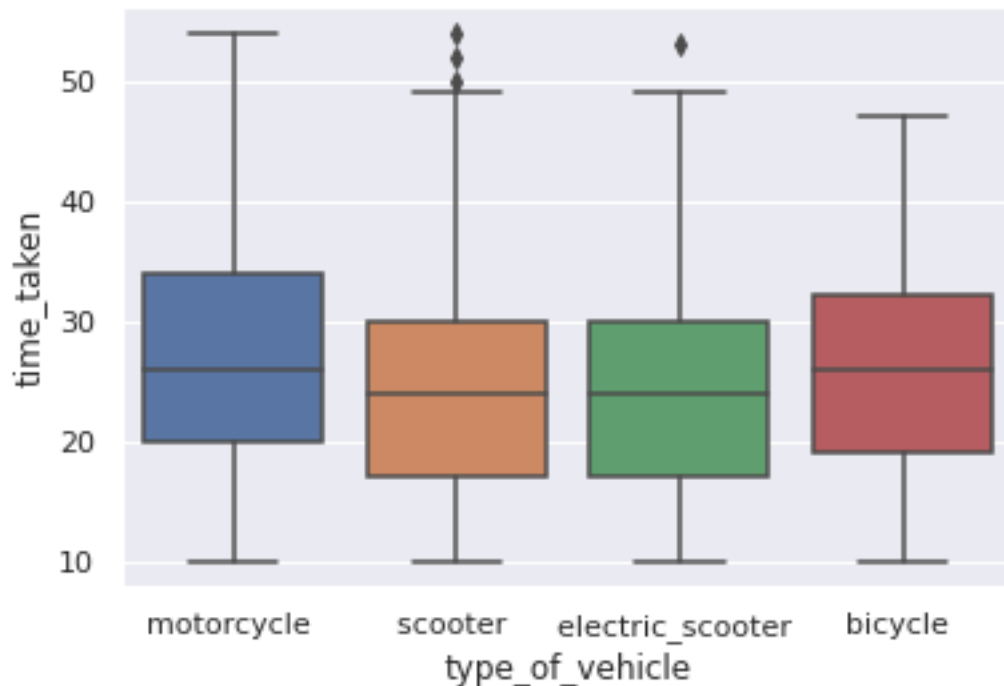
Number of unique vehicle types: 4

```
[255]: ['motorcycle ', 'scooter ', 'electric_scooter ', 'bicycle ']
```

```
[256]: cat_cols["type_of_vehicle"] = [cell.strip() for cell in cat_cols.
↳ type_of_vehicle]
```

```
[257]: sns.boxplot(x = df.type_of_vehicle, y = df.time_taken)
```

```
[257]: <AxesSubplot:xlabel='type_of_vehicle', ylabel='time_taken'>
```



```
[258]: df.groupby("type_of_vehicle")["time_taken"].describe()
```

```
[258]:
```

	count	mean	std	min	25%	50%	75%	max
type_of_vehicle								
bicycle	68.0	26.426471	9.262855	10.0	19.0	26.0	32.25	47.0
electric_scooter	3814.0	24.470110	8.610859	10.0	17.0	24.0	30.00	53.0
motorcycle	26435.0	27.605674	9.647811	10.0	20.0	26.0	34.00	54.0
scooter	15276.0	24.480754	8.704238	10.0	17.0	24.0	30.00	54.0

Deliveries brought by motorcycle or bicycle tend to take longer to deliver. This doesn't necessarily have to be a causal relationship. It could simply be that motorcycles are used to reach destinations that are farther away. It's also of note that only very few deliveries are done by bicycle, so the summary statistics might not be very meaningful in this case.

1.4.3 Delivery person

```
[259]: print(f"Number of unique delivery person ids: {cat_cols.delivery_person_id.
↪nunique()}")
```

Number of unique delivery person ids: 1320

```
[260]: person_id_table = cat_cols.delivery_person_id.value_counts()
person_id_table.describe()
```

```
[260]: count    1320.000000
mean         34.540152
std          21.305850
min           5.000000
25%          13.000000
50%          41.000000
75%          56.000000
max          67.000000
Name: delivery_person_id, dtype: float64
```

The average delivery person has carried out 34.5 deliveries. The delivery person with the fewest number of deliveries has 5, while the most active delivery person has 67 deliveries under their belt. Considering the high number of delivery persons and the relatively few deliveries per person, dummifying this variable might be more trouble than it is worth. It is dropped for now.

```
[261]: cat_cols.drop("delivery_person_id", axis = 1, inplace = True)
```

1.4.4 Dummifying

The type of vehicle and the type of order have relatively few unique categories and they might have a relation to the outcome variable. These two variables will be dummified, while the other two variables (order ID and delivery person ID) have been dropped.

```
[262]: cat_cols_dummy = pd.get_dummies(cat_cols)
print(cat_cols_dummy.columns)
```

```
Index(['type_of_order_Buffet', 'type_of_order_Drinks', 'type_of_order_Meal',
      'type_of_order_Snack', 'type_of_vehicle_bicycle',
      'type_of_vehicle_electric_scooter', 'type_of_vehicle_motorcycle',
      'type_of_vehicle_scooter'],
      dtype='object')
```

For each variable, the first category is dropped and used as reference category. In the case of the type of order, the category I've chosen as reference is “Drinks” as they are easy to prepare and slightly faster to deliver. The type of vehicle chosen as the reference category is “Motorcycle” because they are most frequently used.

```
[263]: cat_cols_dummy.drop(["type_of_order_Drinks", "type_of_vehicle_motorcycle"],
    ↪axis = 1, inplace = True)
cat_cols_dummy.head()
```

```
[263]:
```

	type_of_order_Buffet	type_of_order_Meal	type_of_order_Snack	\
0	0	0	1	
1	0	0	1	
2	0	0	0	
3	1	0	0	
4	0	0	1	

	type_of_vehicle_bicycle	type_of_vehicle_electric_scooter	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	type_of_vehicle_scooter
0	0
1	1
2	0
3	0
4	1

1.5 Numeric independent variables

1.5.1 Distance

The variables ‘restaurant_latitude’, ‘restaurant_longitude’, ‘delivery_location_latitude’, ‘delivery_location_longitude’ can be used to calculate beeline location. Of course, delivery drivers can’t drive in a straight line to their location, but calculating exact distance would be much more difficult. So this is still a good enough approximation.

But first, there seem to be both negative and positive values in the data. Either these locations are very far apart, or there is some error in the data.

```
[264]: df.loc[df.restaurant_latitude < 0, :]
```


[264]:

	id	delivery_person_id	delivery_person_age	delivery_person_ratings	\
92	C042	AGRRES010DEL01	34	4.7	
283	C044	AGRRES12DEL03	32	4.7	
289	4DB	PUNERES02DEL03	29	4.6	
425	C003	DEHRES13DEL02	29	4.6	
534	473	MYSRES07DEL03	29	4.6	

...	
44933	C0C2	AURGRES03DEL03	30	4.9	
45020	56B	PUNERES04DEL01	29	4.6	
45108	C01B	GOARES18DEL01	29	4.6	
45182	C0C7	KNPRES16DEL02	30	4.0	
45504	461	BANGRES07DEL02	29	4.6	

	restaurant_latitude	restaurant_longitude	delivery_location_latitude	\
92	-27.163303	78.057044	27.233303	
283	-27.165108	78.015053	27.225108	
289	-18.551440	-73.804855	18.611440	
425	-30.366322	-78.070453	30.496322	
534	-12.325461	-76.632278	12.385461	
...	
44933	-19.874733	75.353942	19.904733	
45020	-18.514210	73.838429	18.524210	
45108	-15.493950	-73.827423	15.563950	
45182	-26.482581	80.315628	26.532581	
45504	-12.978453	-77.643685	12.998453	

	delivery_location_longitude	type_of_order	type_of_vehicle	\
92	78.127044	Drinks	scooter	
283	78.075053	Meal	scooter	
289	73.864855	Meal	scooter	
425	78.200453	Snack	scooter	
534	76.692278	Buffet	scooter	
...	
44933	75.383942	Buffet	motorcycle	
45020	73.848429	Drinks	electric_scooter	
45108	73.897423	Meal	scooter	
45182	80.365628	Drinks	motorcycle	
45504	77.663685	Meal	scooter	

	time_taken	restaurant
92	15	AGRRES01
283	31	AGRRES12
289	12	PUNERES02
425	20	DEHRES13
534	16	MYSRES07
...
44933	15	AURGRES03

45020	30	PUNERES04
45108	19	GOARES18
45182	34	KNPRES16
45504	24	BANGRES07

[431 rows x 12 columns]

As we can see here, the absolute value of the restaurant latitude seems to be very close to delivery location latitude. As it is highly unlikely that the delivery has crossed the equator or the prime meridian, the values have to be corrected. The locations are likely all in India (for example, [30.496322, 78.200453] is in Uttarakhand), so both positive latitude and positive longitude is assumed.

There are also many restaurants with a location of (0, 0).

```
[265]: tmp = df.loc[df.restaurant_latitude == 0, :]

print(tmp.delivery_location_latitude.describe())
print(tmp.delivery_location_longitude.describe())
```

```
count      3640.000000
mean         0.063016
std          0.036047
min          0.010000
25%          0.030000
50%          0.060000
75%          0.090000
max          0.130000
Name: delivery_location_latitude, dtype: float64
count      3640.000000
mean         0.063016
std          0.036047
min          0.010000
25%          0.030000
50%          0.060000
75%          0.090000
max          0.130000
Name: delivery_location_longitude, dtype: float64
```

In these cases, the delivery location latitude and longitude are all positive and closeby.

The absolute value of all location variables should be taken.

```
[266]: df["restaurant_latitude"] = df["restaurant_latitude"].abs()
df["restaurant_longitude"] = df["restaurant_longitude"].abs()
df["delivery_location_latitude"] = df["delivery_location_latitude"].abs()
df["delivery_location_longitude"] = df["delivery_location_longitude"].abs()
```

Finally, we can calculate the beeline distance.

```
[267]: def calculate_distance(lat1, lat2, lon1, lon2):
        """ Calculate the distance between two points using the haversine formula """
        ↪
        lat1, lat2, lon1, lon2 = map(np.radians, [lat1, lat2, lon1, lon2])

        dlat = lat2 - lat1
        dlon = lon2 - lon1
        a = np.sin(dlat/2) ** 2 + np.cos(lat1) * np.cos(lat2) * np.sin(dlon/2) ** 2
        b = 2 * np.arctan2(a ** 0.5, (1-a) ** 0.5)

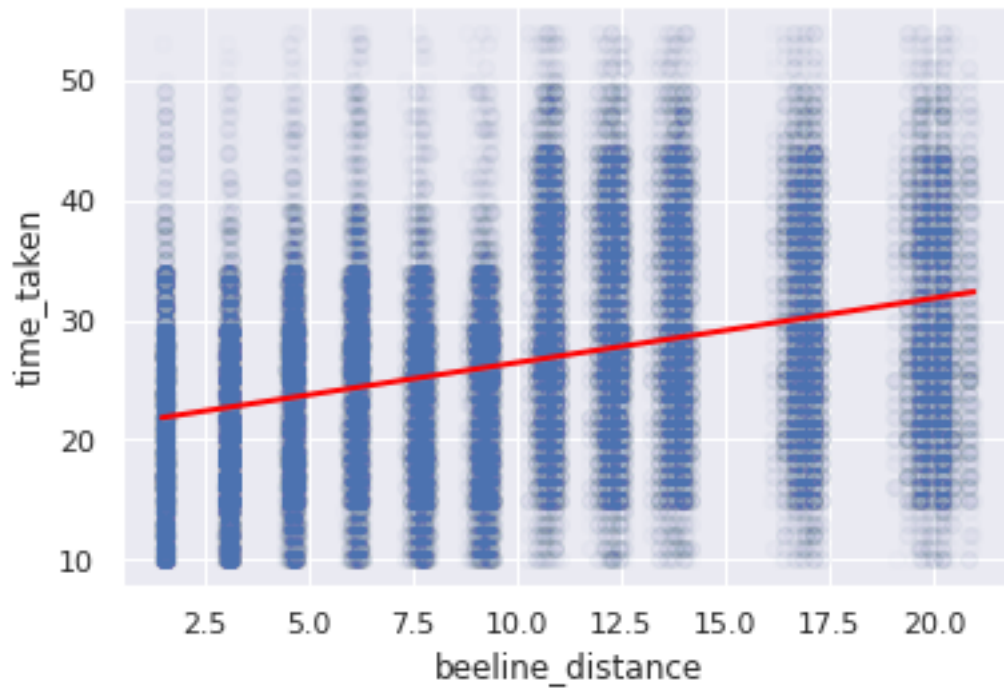
        return b * 6371

df["beeline_distance"] = calculate_distance(df.restaurant_latitude, df.
        ↪delivery_location_latitude,
                                           df.restaurant_longitude, df.
        ↪delivery_location_longitude)
beeline_distance.describe()
```

```
[267]: count    45593.000000
        mean       0.907852
        std       10.087993
        min       0.014141
        25%       0.042426
        50%       0.084853
        75%       0.127279
        max       182.546210
        dtype: float64
```

```
[268]: sns.regplot(x = df.beeline_distance, y = df.time_taken,
                    scatter_kws = {'alpha':0.01},
                    line_kws = {"color": 'red', "linewidth": 2})
```

```
[268]: <AxesSubplot:xlabel='beeline_distance', ylabel='time_taken'>
```

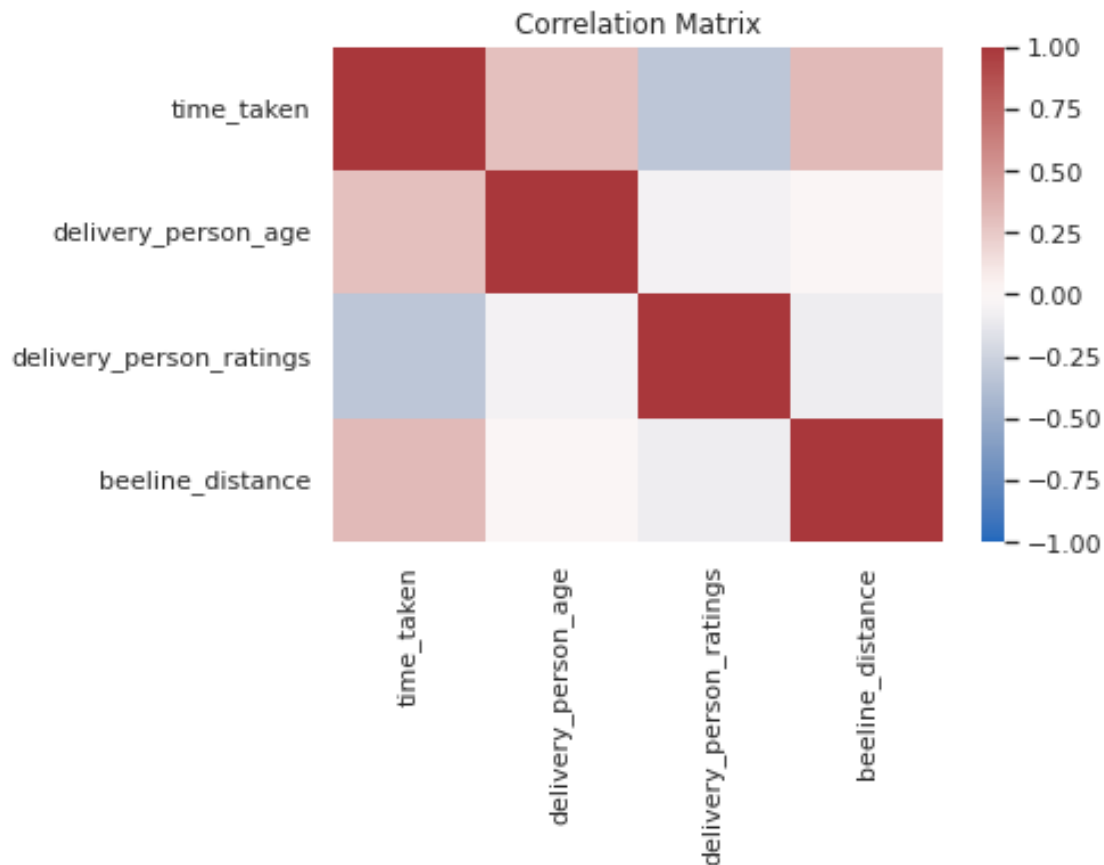


There seems to be a positive association between the beeline distance and the time taken.

1.5.2 Age and rating of delivery person

```
[269]: df_num = df[["time_taken", "delivery_person_age", "delivery_person_ratings",
↪ "beeline_distance"]]
corr_matrix = df_num.corr()
heatplot = sns.heatmap(corr_matrix, cmap = "vlag", vmin = -1, vmax = 1)
heatplot.set(title = "Correlation Matrix")
```

```
[269]: [Text(0.5, 1.0, 'Correlation Matrix')]
```



- Older delivery persons take more time for the delivery
- Delivery persons with high ratings deliver more quickly
- The associations between the independent variables is low, so we expect no issues with multicollinearity

```
[270]: print(df.delivery_person_ratings.describe())
print(f"Number of cases with a rating of 6: {sum(df.delivery_person_ratings == 6)}")
```

```
count    45593.000000
mean      4.632367
std       0.327708
min       1.000000
25%      4.600000
50%      4.700000
75%      4.800000
max       6.000000
Name: delivery_person_ratings, dtype: float64
Number of cases with a rating of 6: 53
```

There are 53 cases with a rating of 6. A rating scale from 1 to 5 is uncommon, and the cases are

very few, so this is probably a mistake. The rating is recoded to 5.

```
[271]: df.loc[df.delivery_person_ratings == 6, "delivery_person_ratings"] = 5
```

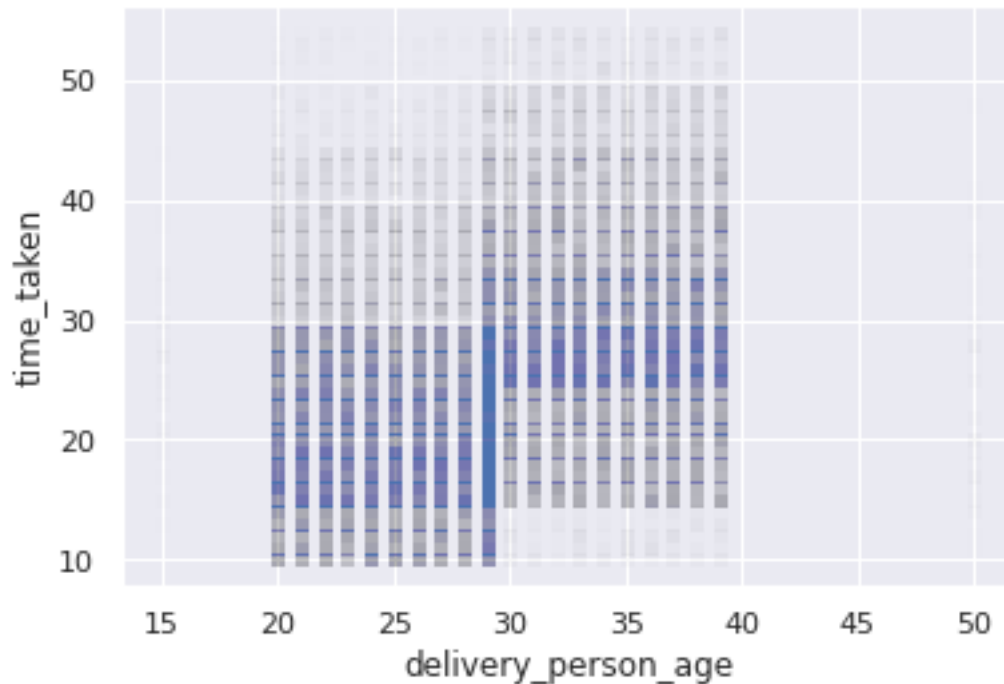
```
[272]: sns.scatterplot(x = df.delivery_person_ratings, y = df.time_taken, alpha = 0.  
↪005)
```

```
[272]: <AxesSubplot:xlabel='delivery_person_ratings', ylabel='time_taken'>
```



```
[273]: sns.scatterplot(x = df.delivery_person_age, y = df.time_taken, alpha = 0.005)
```

```
[273]: <AxesSubplot:xlabel='delivery_person_age', ylabel='time_taken'>
```



1.6 Save dataset

```
[274]: print(df.columns)
       print(cat_cols_dummy.columns)
```

```
Index(['id', 'delivery_person_id', 'delivery_person_age',
      'delivery_person_ratings', 'restaurant_latitude',
      'restaurant_longitude', 'delivery_location_latitude',
      'delivery_location_longitude', 'type_of_order', 'type_of_vehicle',
      'time_taken', 'restaurant', 'beeline_distance'],
      dtype='object')
Index(['type_of_order_Buffet', 'type_of_order_Meal', 'type_of_order_Snack',
      'type_of_vehicle_bicycle', 'type_of_vehicle_electric_scooter',
      'type_of_vehicle_scooter'],
      dtype='object')
```

```
[275]: df_finished = pd.concat([
        df[["time_taken", "delivery_person_age", "delivery_person_ratings",
        ↪ "beeline_distance"]],
        cat_cols_dummy
    ], axis = 1)
df_finished.head()
```

```

[275]:  time_taken  delivery_person_age  delivery_person_ratings  beeline_distance  \
0          24             37             4.9             3.025149
1          33             34             4.5            20.183530
2          26             23             4.4             1.552758
3          21             38             4.7             7.790401
4          30             32             4.6             6.210138

      type_of_order_Buffet  type_of_order_Meal  type_of_order_Snack  \
0                      0                0                1
1                      0                0                1
2                      0                0                0
3                      1                0                0
4                      0                0                1

      type_of_vehicle_bicycle  type_of_vehicle_electric_scooter  \
0                      0                0
1                      0                0
2                      0                0
3                      0                0
4                      0                0

      type_of_vehicle_scooter
0                      0
1                      1
2                      0
3                      0
4                      1

```

```

[276]: df_finished.to_csv("data/deliverytime_processed.csv")

```